

ASSIGNMENT № 3

Arun K Rajagopalan, Texas A&M University

11/09/2015

Problem

Apply MapReduce to graph relabeling and shows the benefits of doing so.

Analysis

The problem can be split into 3 stages -

1. Graph Inversion 1
2. Graph Inversion 2
3. Compute TopN

Each of these stages have been designed to operate independently of the previous stage. By treating each stage as a black-box, I was able to efficiently debug and ensure correctness.

Graph Inversion

Algorithm

I present an overview of the Algorithm at a high level.

- The graph inversion step 1 and step 2 are basically the same, with minor modifications. Thus, I templated the graph inversion. It takes as template an argument that specifies the size of the vertex node (either uint64 or uint32). Specializations are provided by means of a boolean flag to differentiate the two stage.
- The inversion proceeds under the assumption that the nodes of the graph arrive in a sorted order.
- We read the nodes and relabel either from the previous run or sequentially depending on whether it is the first run or not.
- For each node we relabel, we map it into a buffer inverted.
- The buffer performs sort-compact.
- We then write the buffer to disk.
- We then need to merge the split files. This is similar to a merge of two sorted arrays.
- Finally, we end up with an output file that is inverted.

Optimizations

There are multiple avenues of optimization that I looked into.

- While mapping the inverted nodes, I send it into 8 buckets, so that I can do a parallel sort. In the first phase, I use the top 3 bits for this. In the second inversion, I cannot do this since the nodes are no longer properly distributed. Instead, I partition on the range of the input.
- The merge is performed as a 4-way merge executed in 2 threads (effectively, 8 files are being merged concurrently).
- To maintain the file handle count, I use a fileHandle map inside each thread that keeps the total count to 2 per thread plus one shared write handle between the two threads.
- I used un-buffered reads to get maximum read performance.

Areas of future improvement

While I was able to achieve good speedup with the above mentioned techniques, these were some of the areas I would have also liked to improve.

- Asynchronous reads will give me much better performance.
- The adjacency list sort in the final step can be distributed to a thread pool. I got a prototype working, but it was consuming too much memory.
- I am missing a few nodes from the final count. This could possibly be due to some data getting cut off at a buffer overrun.

TopN

Algorithm

I present an overview of the Algorithm at a high level.

- Since the relabelled graph now fits in memory, we can directly increment the counter for each node as we see them.
- Once this is done, we create a priority queue of size 10. This queue holds the top 10 elements in the array as we scan it.

Non-trivial and interesting points

- memmove should be used instead of memcpy if the dst buffer overlaps with source.

Statistics

Following is the output when run on my personal Macbook Pro (Intel Core i5-4278@2.6GHz, 8GB RAM, SSD) running Windows 10.

```
Starting 1st Inversion
Total IO: read - 53.40 GB; write - 47.13 GB
Took 252 seconds
Ending 1st Inversion
```

```
Starting 2nd Inversion
Total IO: read - 39.64 GB; write - 38.60 GB
Took 272 seconds
Ending 2nd Inversion
```

```
Overall stats - RunTime: 525 seconds;
Total read 93.03 GB; Total write 85.73 GB
Peak working set : 499.70 MB
DONE!
```

```
-----
```

```
Starting Top 10
```

```
TOP 10 Node count
```

```
1. 24247151 2947626
2. 21153020 2209833
3. 17991418 1996682
4. 13620962 1287415
5. 14499776 1203819
6. 3229951 1031992
7. 173482 932753
8. 14715032 932356
9. 9904332 803930
10. 18092264 744919
```

```
Total sum  $Y_i(Y_{i-1})/2 = 3144604790303$ 
Total IO: read - 7.41 GB; write - 0.00 GB
Took 53 seconds
Ending Top 10
```

Overall stats - RunTime: 53 seconds;
Total read 7.41 GB; Total write 0.00 GB
Peak working set : 459.32 MB
DONE!

Top 10 links

1 microsoft.com 2947630
2 google.com 2209834
3 yahoo.com 1996682
4 adobe.com 1287417
5 blogspot.com 1203819
6 wikipedia.org 1031992
7 geocities.com 932753
8 w3.org 932359
9 msn.com 803930
10 amazon.com 744919

ATTRIBUTION

I took help from the following sites to help with my assignment

1. stackoverflow.com
2. MSDN
3. sparetimelabs.com/printfrevisited